

Convolutional Neural Network Architecture in NumPy

Aaron Avram

June 14 2025

Introduction

In this writeup I am going to outline the design of my Convolutional Neural Network for image classification. However, my hope is to build a module that abstracts away individual model design and provides a framework for deploying any model type, drawing some inspiration from the architecture of PyTorch.

High-Level Design

At the foundation of my model module is the Tensor class. This is a class with two NumPy arrays as attributes, one representing a value matrix and the other a gradient matrix. Since practically all of the components of a Neural network store value and gradient data it felt natural to couple them under a single class. The Tensor class will support gradient update and setting the gradient to zero.

On top of this class is the Layer class. This is meant to capture the core functionality of a layer in a neural network: it has an input, an (optional) list of parameters and an output, all stored as Tensor objects. The class supports a forward and a backward pass. I will create subclasses of the layer class for individual layer types, e.g. Convolutional, Dense or Activation. However, the superclass allows for polymorphic behaviour when implementing an entire model.

Finally, I am going to build the model class which has a list of layers, which chain together from input to output, using aliasing to conserve memory. My CNN model will be a particular instance of the model class with its layers structured in the standard way for image classifying CNNs.

Mathematical details

I will not go into the details of backpropagation for the dense layers in my model, as I have already derived these things in my NameNet project and will instead focus on the Convolutional, Batch Norm, and Pooling layers of my model.

Convolutions

The first layer in my architecture takes in a batch of training examples $X^{(in)}$ which has dimensions (n, c_{in}, d, d) , assuming square input data. Where n is the size of the batch. c_{in} is the number of channels in the input data, and d is the size of each training point. We then convolve this tensor with a kernel tensor K of dimension (c_{out}, c_{in}, k, k) , where c_{out} is the number of points in the next

layer and K is the size of the kernel tensor. Now the Convolution operation on these tensors is defined by:

$$(X^{(in)} * K)_{i,j,k,l} = \sum_{r,s,t} X_{i,r,k+s,l+t}^{(in)} K_{r,s,t}$$

This is frankly a bit of a mess and inefficient as its time complexity scales linearly in each argument of $X^{(in)}$ and K getting out of hand quickly.

Luckily there is a more efficient way to perform this operation, as we can see that K slides along the each axis of X acting as a sort of sighting window. Thus, we can first copy out the necessary part of X then implement this operation as a matrix operation between the reshaped components. For this we will use a function provided by *NumPy* to make a copy of X and we will call this $X^{(col)}$, and if we let p be the size of the output points we see that $X^{(col)}$ will have dimension (n, c_{in}, p, p, k, k) . So $X^{(col)}$ has all of the windows stored. Next we reorder the indices of $X^{(col)}$ with NumPy's transpose function to give it dimension n, p, p, c_{in}, k, k then we reshape it to have dimension $(np^2, c_{in}k^2)$. Next we reshape K to dimension $(c_{out}, c_{in}k^2)$. Then we can calculate their convolution by taking:

$$X^{(out)} = X^{(col)} K^T$$

And if we reshape $X^{(out)}$ we can give it the dimension n, c_{out}, p, p . This implementation is orders of magnitude faster than the previous implementation.

A note, often times these convolutions are done using strides, i.e. skipping certain windows when constructing $X^{(col)}$, this doesn't really change any derivations above or further down, only burdens notation.

Convolution Gradient